

VM BASED TRUSTED EXECUTION ENVIRONMENTS ON KUBERNETES

Project **Raksh** (*Protect*)

<https://github.com/ibm/raksh>

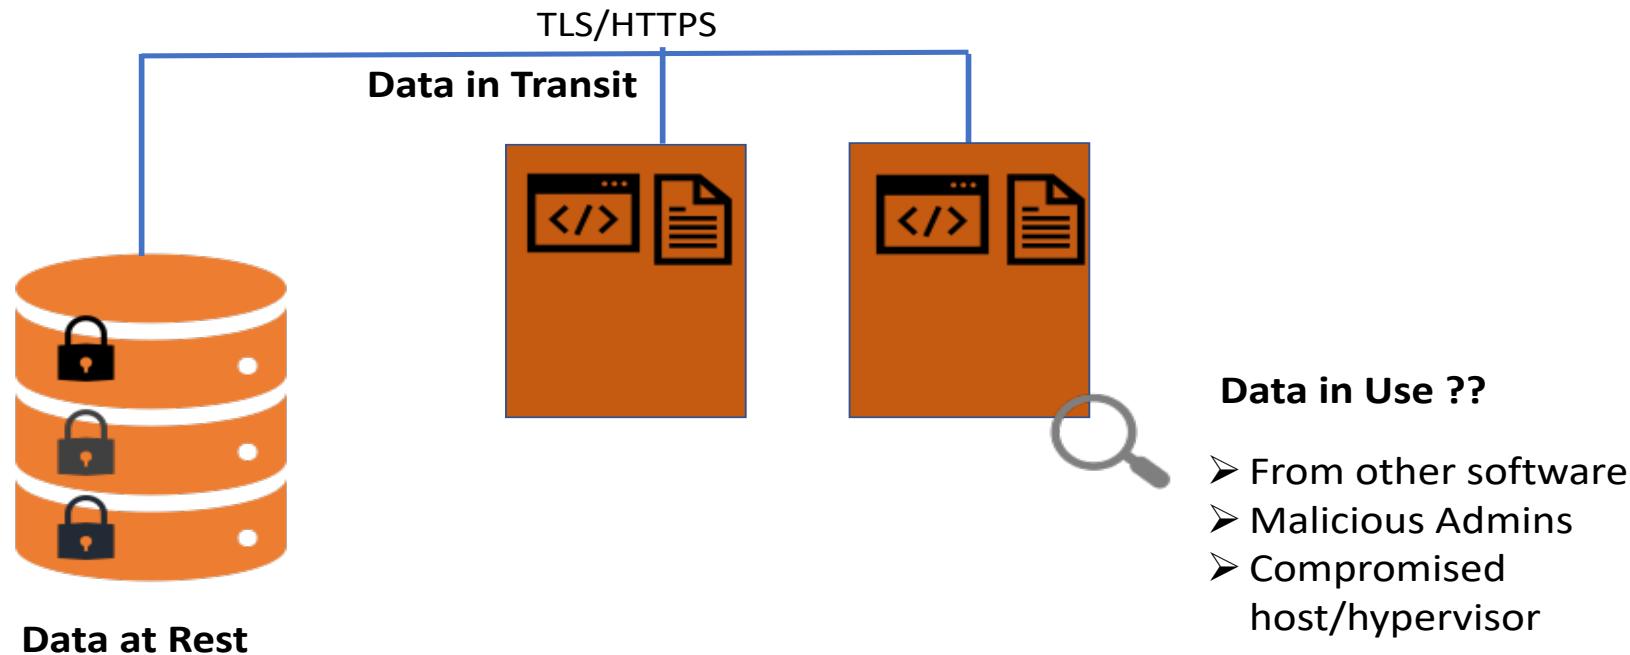


AGENDA

- Problem
- Trusted Execution Environment
- Raksh Overview
- Next Steps



PROBLEM – PROTECT DATA IN-USE

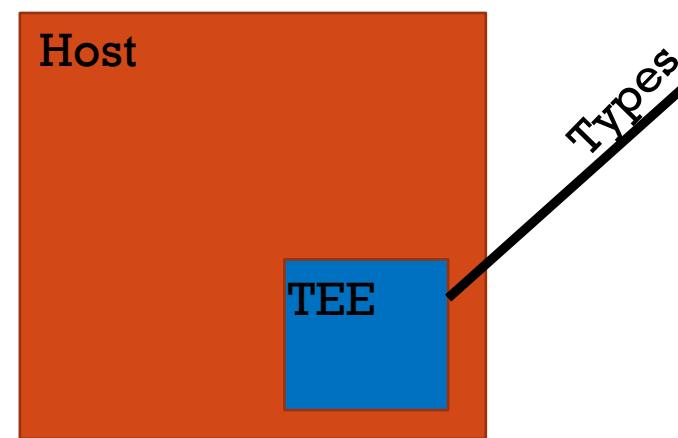


How do you maintain data confidentiality when running sensitive workloads on your premises or in the cloud ?



SOLUTION - TRUSTED EXECUTION ENVIRONMENT (TEE)

- TEE is a protected area within the host used for execution of sensitive workloads
- TEEs provide confidentiality guarantees for code and data running within them. IOW the running workload can't be seen from outside the TEE.



- **Process based**
 - Process address space obfuscated from the OS
 - Examples - Intel SGX
 - Application code modification required
 - No virtualization overheads.
- **VM based**
 - VM address space obfuscated from the hypervisor
 - Examples – Intel TME/MKTME, AMD SEV, IBM Power9 PEF
 - No modifications to application code required
 - Virtualization overhead





INDUSTRY NEED

How to use TEEs
with Kubernetes
(Openshift, IKS,
AKS, PKS ...)?

PROJECT RAKSHI

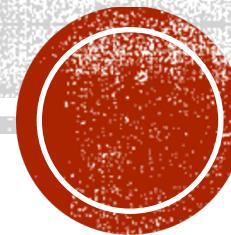


INTRODUCTION

- **Salient Points**
 - **Raksh** - means Protect
(<https://github.com/ibm/raksh>)
 - Opensource project to enable deployment of **VM based TEEs** in a Kubernetes cluster
 - Introduces **Secure Containers** resource in K8s which are **containers protected by VM based TEEs**
 - Introduces **encryption for K8s application spec** (pod.yaml, deployment.yaml etc)
 - No changes to K8s application deployment workflow
 - Leverages Kata VM container runtime
 - Built on Kubernetes Operator pattern
 - CLI tool (*rakshctl*)
- **Dependencies**
 - Supported Hardware
 - Kata container runtime



RAKSH DETAILS



SECURE CONTAINERS

- **Raksh** defines **Secure container** as a container that runs in VM based TEE. The container is isolated and protected using VM based TEE.
- **Secure Container in Raksh** is implemented using the **Operator** pattern.
- Following are the Custom Resource Definitions (CRDs)
 - **SecureContainer** - Manages the lifecycle of the Secure Container
 - **SecureContainerImage** - Manages the lifecycle of VM image used to execute the Secure Container
 - **SecureContainerImageConfig** - Manages the configuration (download location etc) of the SecureContainerImage



SECURE CONTAINER CRDS

```
apiVersion: securecontainers.k8s.io/v1alpha1
kind: SecureContainerImageConfig
metadata:
  name: <config-resource-name>
spec:
  imageDir: <host-path-to-store-image>
  runtimeClassName: kata-containers
```

SecureContainerImageConfig

- Kata VM kernel and initrd location
- Kata Runtime Class to use

```
apiVersion: securecontainers.k8s.io/v1alpha1
kind: SecureContainerImage
metadata:
  name: <image-resource-name>
spec:
  vmImage: <container-image-name>:<tag>
  imagePullSecrets:
    - name: <registry-secret>
  imagePullPolicy: Always
  SecureContainerImageConfigRef:
    name: <config-resource-name>
```

SecureContainerImage

- Container image with Kata VM kernel and initrd
- Download and Extract to location specified by SecureContainerImageConfig

```
apiVersion: securecontainers.k8s.io/v1alpha1
kind: SecureContainer
metadata:
  name: <securecontainer-resource-name>
object:
  apiVersion: v1
  kind: Pod
  metadata:
    name: <app-name>
  spec:
    containers:
      <...>
    volumes:
      <...>
  spec:
    SecureContainerImageRef:
      name: < image-resource-name>
```

SecureContainer

- Secure Container resource to deploy app container protected using VM based TEE



CUSTOM KATA AGENT

- Handle container image lifecycle (download, extract, decryption) inside the VM
- Execute the containers inside VM



RAKSH CLI

- **rakshctl**
 - CLI to convert an existing Kubernetes application specification (pod.yaml, deployment.yaml etc) to support deployment leveraging VM based TEE



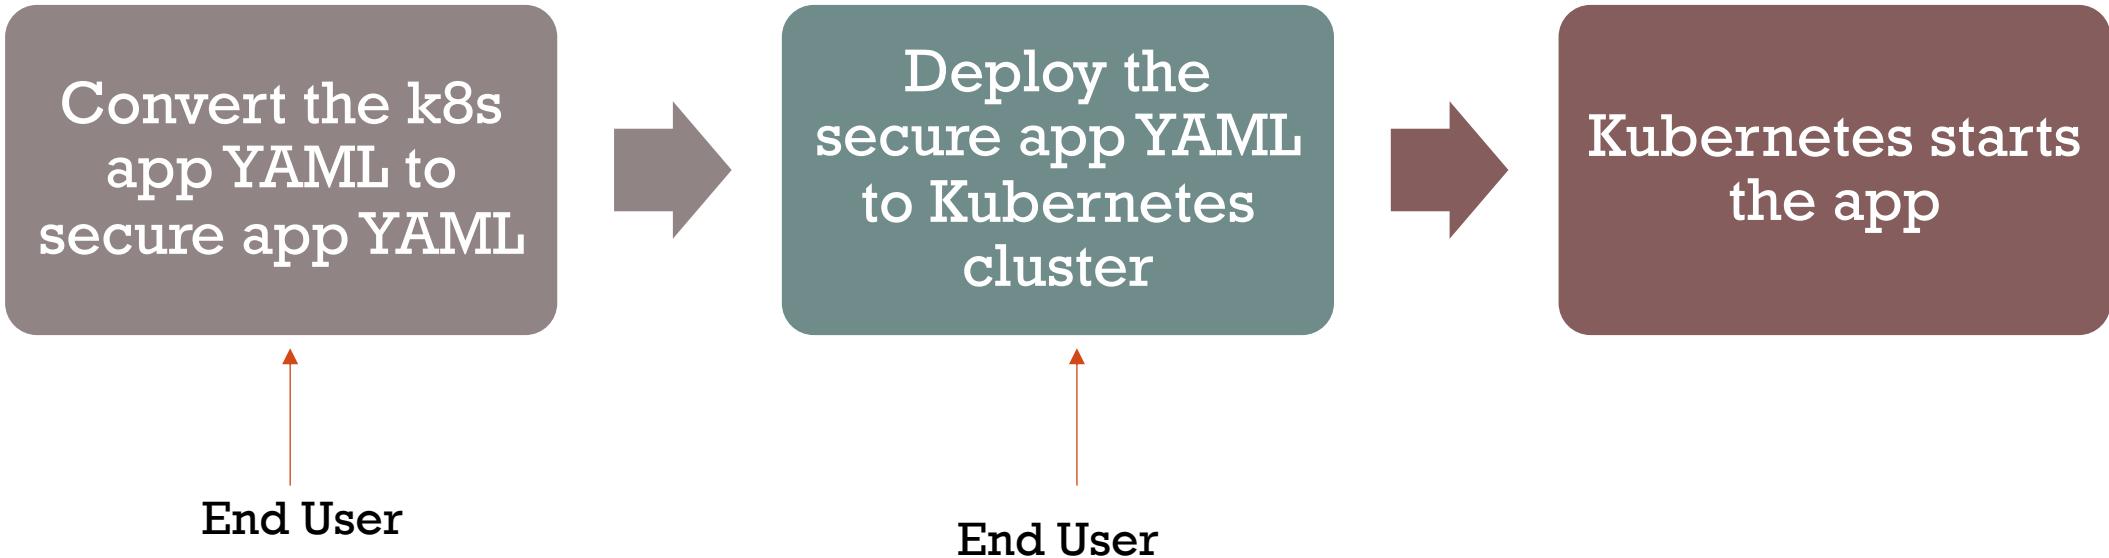
APPLICATION DEPLOYMENT WORKFLOW

- **User Story**

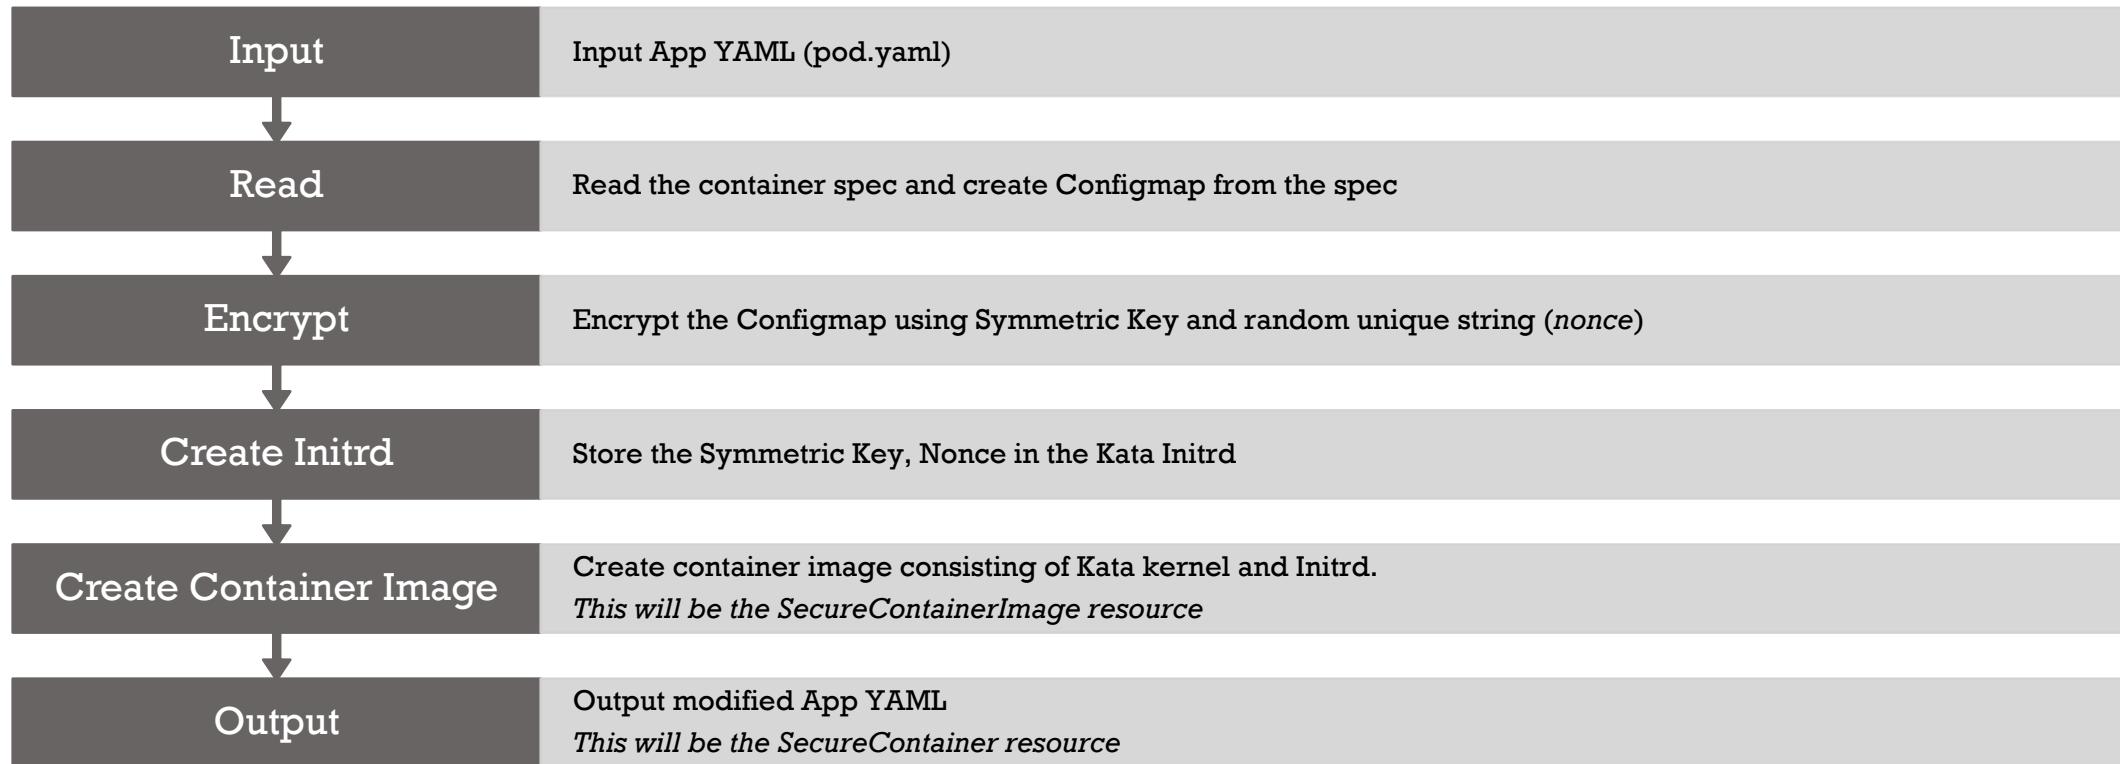
- As a user I want to deploy and run my Kubernetes application securely with minimal disruption to my workflow and protect the data in-use from malicious entities



STEPS



GENERATE SECURE APP YAML



GENERATE SECURE APP YAML

Convert the YAML to secure YAML

```
rakshctl image create -i nginx-securecontainerimage --initrd kata-containers-initrd.img --symmKey /home/pradipta/symm_key --vmlinux vmlinuz --scratch sc-scratch:latest --filename /home/pradipta/nginx.yaml nginx-securecontainerimage
```

1. Create a ConfigMap of the container spec

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    app: nginx
  namespace: default
  name: nginx
spec:
  containers:
    - image: nginx:latest
      imagePullPolicy: IfNotPresent
      name: nginx
      ports:
        - containerPort: 80
          protocol: TCP
```

nginx.yaml

1

```
apiVersion: v1
data:
  nginx: |
spec:
  containers:
    - image: nginx:latest
      imagePullPolicy: IfNotPresent
      name: nginx
      ports:
        - containerPort: 80
          protocol: TCP
  kind: ConfigMap
  metadata:
    name: configmap-nginx
```

ConfigMap

CONT...

Convert the YAML to secure YAML

```
rakshctl image create -i nginx-securecontainerimage --initrd kata-containers-initrd.img --symmKey /home/pradipta/symm_key --vmlinuz vmlinuz --scratch sc-scratch:latest --filename /home/pradipta/nginx.yaml nginx-securecontainerimage
```

1. Create a ConfigMap of the container spec
2. Encrypt the ConfigMap with symmetric key

```
apiVersion: v1
data:
  nginx: |
    spec:
      containers:
        - image: nginx:latest
          imagePullPolicy: IfNotPresent
          name: nginx
          ports:
            - containerPort: 80
              protocol: TCP
kind: ConfigMap
metadata:
  name: configmap-nginx
```

2

```
apiVersion: v1
data:
  nginx:
    6qvygg8md7bXfyX3Y9cpZxUp4eZA0kKmWBirrpJv/WEGkrdLYrdtqxdqm4cGL
    G4++06d2iGTaB+5SDjjDwf05T+9a2iUAdHmRngHcQNAzkKK2RCnR4Zkt0cXDa
    EP+w5mbugH0xdqGm8SoX4IgvWGi2toq1CUcc8OmgTX42g0NruTZbrNv5Ncc
    yS7+kR7lib6vaMI24E=
kind: ConfigMap
metadata:
  name: secure-configmap-nginx
```

Encrypted ConfigMap

ConfigMap



CONT...

Convert the YAML to secure YAML

```
rakshctl image create -i nginx-securecontainerimage --initrd kata-containers-initrd.img --symmKey /home/pradipta/symm_key --vmlinuz vmlinuz --scratch sc-scratch:latest --filename /home/pradipta/nginx.yaml nginx-securecontainerimage
```

1. Create a ConfigMap of the container spec
2. Encrypt the ConfigMap with symmetric key (key either added to lockbox or vault)
3. Create a modified container spec to use encrypted ConfigMap (uses dummy image which just sleeps. Actual container image will be pulled inside the Kata VM)

```
apiVersion: v1
data:
  nginx:
    6qvygg8md7bXfyX3Y9cpZxUp4eZA0kKmWBirrpJv/WEGkrdLYrdtqxdqm4cGLG4++06d2iGTaB+5SDjjDwf05T+9a2iUAdHmRngHcQNAzkKK2RCnR4Zkt0cXDaEP+w5mbugH0xdqGm8SoX4IgvWGi2toq1CUcc8OmgTX42g0NruTZbrNv5NccyS7+kR7Iib6vaMI24E=
kind: ConfigMap
metadata:
  name: secure-configmap-nginx
```

3

```
apiVersion: securecontainers.k8s.io/v1alpha1
kind: SecureContainer
metadata:
  name: secure-nginx
object:
  apiVersion: v1
  kind: Pod
  metadata:
    labels:
      app: nginx
      name: nginx
  spec:
    containers:
      - image: sc-scratch:latest
        imagePullPolicy: IfNotPresent
        name: nginx
        ports:
          - containerPort: 80
            protocol: TCP
        resources: {}
        volumeMounts:
          - mountPath: /etc/raksh
            name: secure-volume-nginx
            readOnly: true
        volumes:
          - configMap:
              items:
                - key: nginx
                  path: kavach.properties
              name: secure-configmap-nginx
              name: secure-volume-nginx
        spec:
          SecureContainerImageRef:
            name: nginx-securecontainerimage
```

Dummy Image (sleep) running on host

Mount the encrypted ConfigMap inside the Kata VM

Reference to the encrypted ConfigMap

Kata VM image used

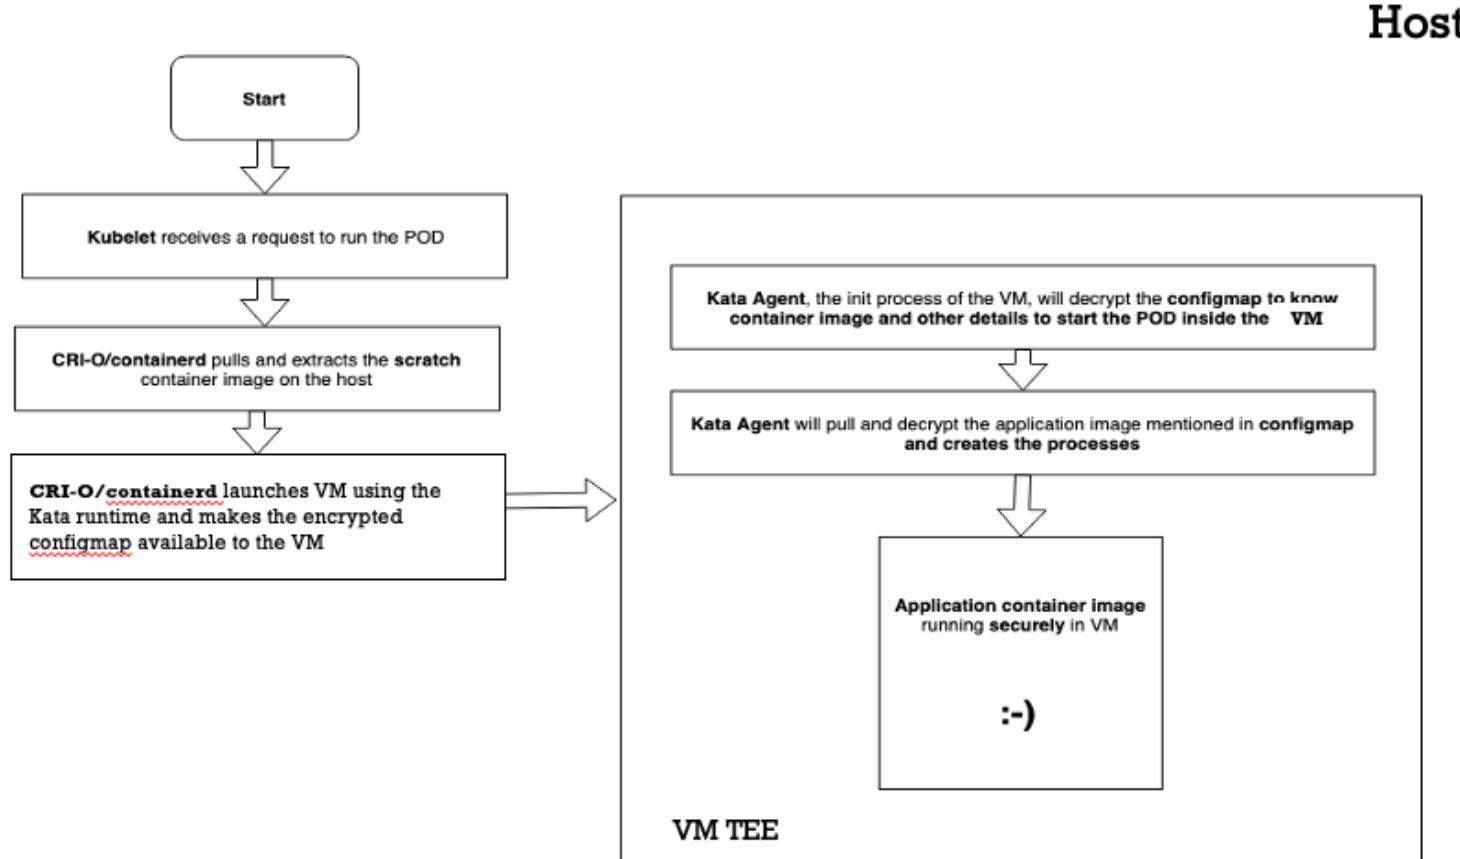
nginx-sc.yaml

DEPLOY THE WORKLOAD

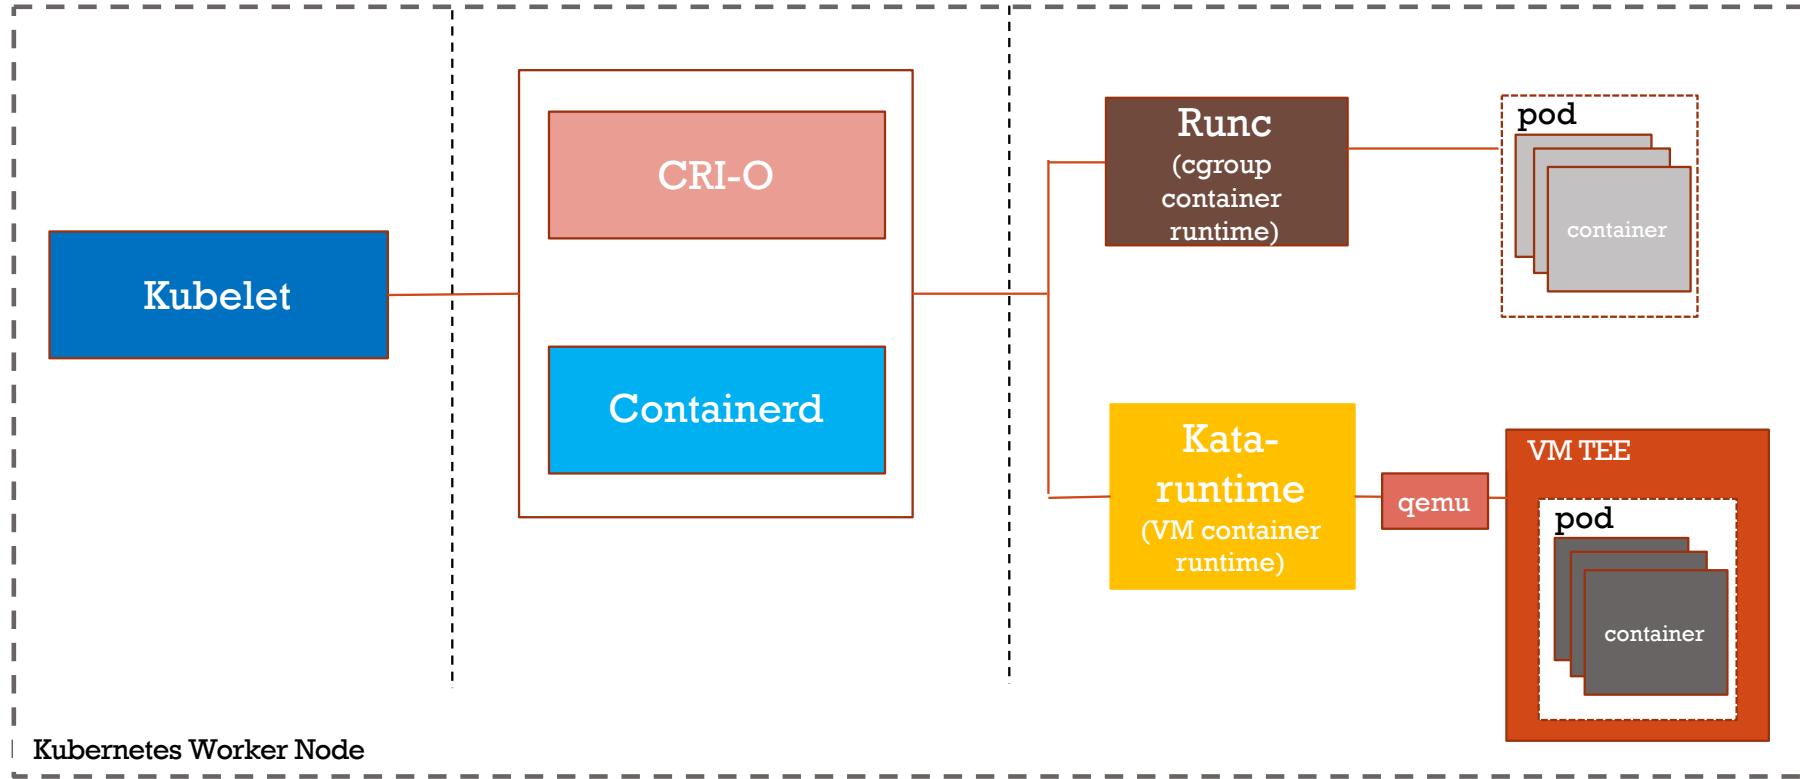
Deploy the Workload

```
kubectl create -f nginx-sc.yaml
```

- All operations w.r.to customer workload is handled inside the Kata VM which is the VM based TEE (AMD SEV, IBM PEF etc)
- The secure workload spec is immutable and it's not possible for anyone with admin privileges to the Kubernetes cluster to modify the workload spec.



COMPONENTS



VM TEE – AMD SEV, IBM PEF, INTEL MKTME



DEMO

1. Create VM image (Kata initrd and vmlinu)
2. Create SecureContainerImageConfig CRD
 1. VM image name, install location etc.
3. Create SecureContainerImage CRD
 1. Docker image with VM kernel, initrd
4. Show vault details stored as Kubernetes secret
5. Show regular app yaml and secure app yaml
6. Deploy the secure app yaml

Demo Link - <https://asciinema.org/a/290233>

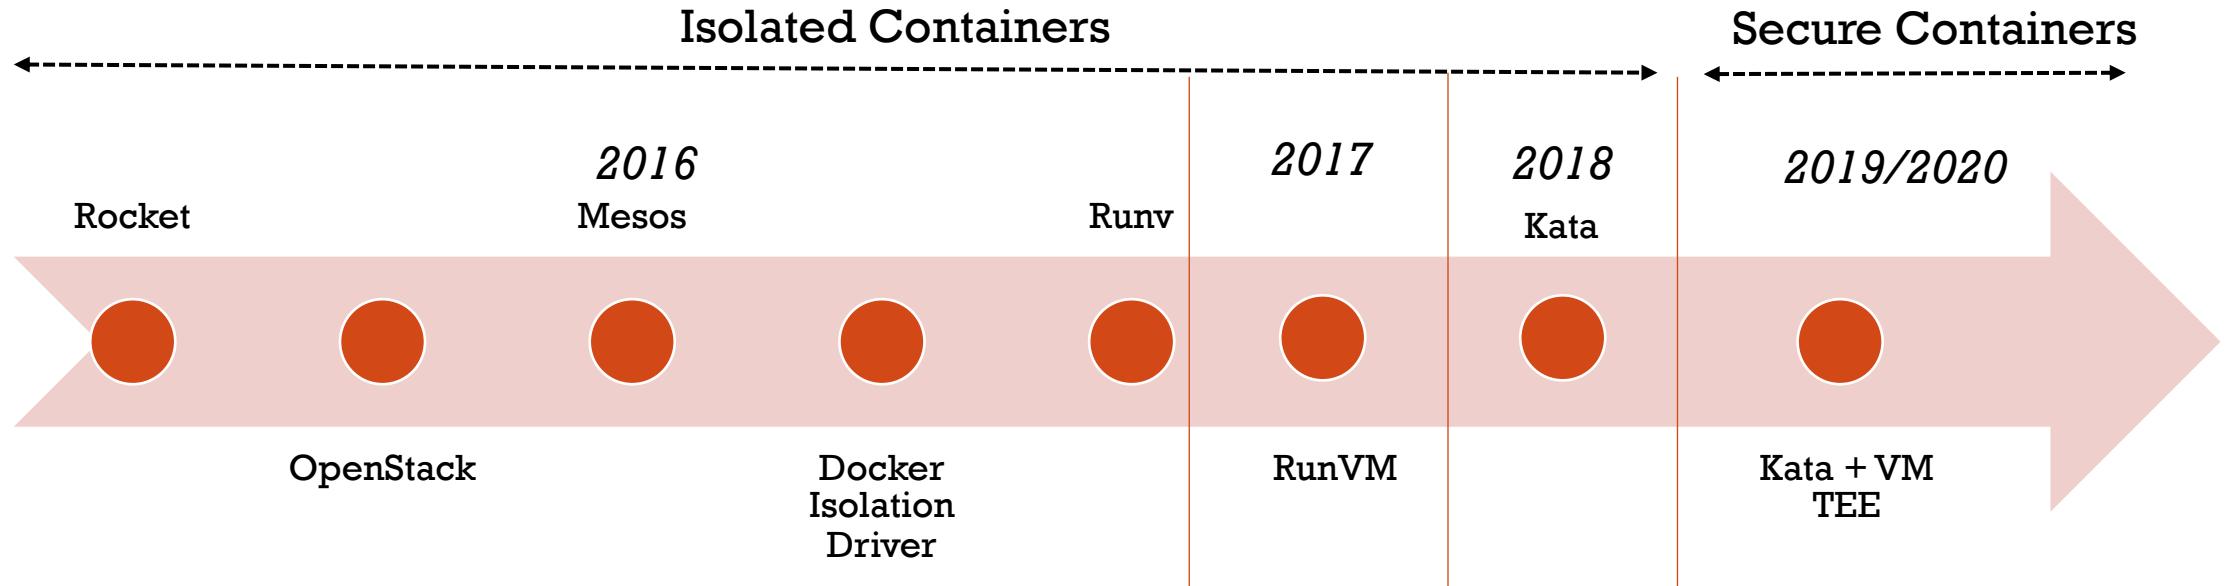


REFERENCES

- Introducing Project Raksh - http://bit.ly/project_raksh
- What You Define is What You Deploy - <http://bit.ly/wydiwud>
- Raksh Github Link - <https://github.com/ibm/raksh>



JOURNEY TO SECURE CONTAINERS



References

- <http://bit.ly/isolatedcontainers>
- <https://github.com/harche/runvm>
- <https://github.com/moby/moby/issues/29454>
- <https://katacontainers.io/posts/kata-containers-ibm-power/>
- <https://github.com/opencontainers/image-spec/issues/747>
- <https://github.com/kata-containers/runtime/issues/61>
- <https://docs.google.com/document/d/1NjM58FZFRLtK2qh90WYnk-1nd8ili1d433d8pGwri6g/edit#heading=h.pext71slr0u9>
- <https://docs.google.com/document/d/1zSk397hmaRuN6LkXzjYjljkxG0SZRw5pr4ZYzfxC4c/edit#heading=h.nupw5ka6qe7w>
- <https://github.com/ibm/raksh>

